

General Information

1. Installation

Both Dask and Spark can be installed with conda

```
conda install dask
```

```
conda install pyspark
```

For Spark, installing Java and setting JAVA_HOME are required

```
sudo apt install openjdk-8-jdk
```

```
sudo nano /etc/environment
```

```
JAVA_HOME="/usr/lib/jvm/java-8-openjdk-amd64"
```

2. Characters

Spark: written in Scala, compatible with Java, Python and R. In general Spark has more ready-to-use tools, e.g. MLlib for machine learning and GraphX for graph processing. It integrates well with many other Apache projects. It is fundamentally an extension of the Map-Shuffle-Reduce paradigm.

Dask: written in Python. It works well with common python libraries like NumPy, Pandas. It is lack of high-level optimization (comparing to Spark), but it is more flexible. Therefore it can build more complex bespoke systems. It is fundamentally based on generic task scheduling.

3. The way of working

- Spark: Spark separate the input parameters into several blocks, and apply parallel computation of all blocks. Within each block the computations are performed sequentially.
 1. Set up SparkContext

```
sc = SparkContext(master="local[{}]" .format(nr_worker))
```
 2. Parallelize input parameters into blocks

```
task = sc.parallelize(in_out_file_pairs_spark)
```
 3. Map the function to each block and collect results

```
task.map(export_ndvi).collect()
```
 4. Stop SparkContext

```
SparkContext.stop(sc)
```
- Dask: Dask registers the operations in a list ("futures" in the script), and distribute these tasks into the available resources defined in the Client.
 1. Set up the client:

```
cluster = LocalCluster(processes=True, n_workers=nr_worker,  
threads_per_worker=1, local_directory='./dask-worker-space')
```
 2. submit the function to a list of futures

```
future = [client.submit(export_ndvi, f) for f in in_out_file_pairs_dask]
```
 3. gather futures (computation)

```
results = client.gather(futures)
```
 4. shutdown client

```
address = client.scheduler.address # get address  
client.close()  
Client(address).shutdown()
```

Test case: NDVI computation with Sentinel-2 data

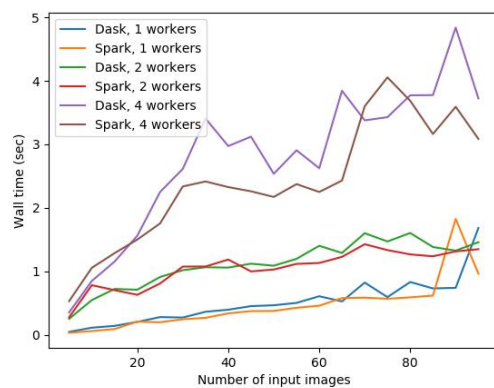
1. Description of the test case

In this test case we attempt to compute the Normalized Difference Vegetation Index (NDVI) from Sentinel-2 images. We run the same NDVI computation through both Spark and Dask, with the same input images, and assess the wall time of computation.

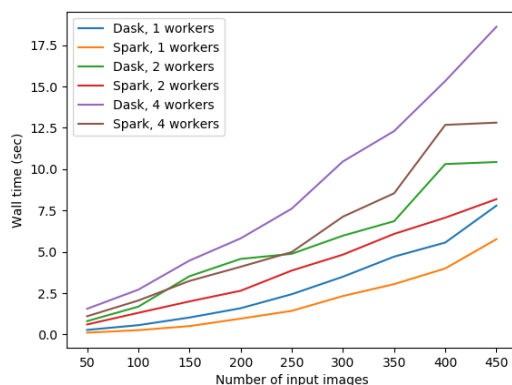
The purpose is to 1) have a local test run and get familiar with both frameworks. 2) investigate if Spark and Dask have major performance difference.

The code of this test can be found in: https://github.com/phenology/dask_spark_comparision_s2

2. Performance comparison



Wall time 0-100 images



Wall time 50-500 images

3. Taking out from the test case

- The performance of Spark and Dask do not show major difference with small amount of images. But when the number of images growing, Spark seem to be slightly faster.
- Strange to see the more workers we have, the slower the computation. Need to check if the number of workers is set correctly. Or if the time is recorded correctly.

Recommendations

Below we assume a Python development case, so we do not consider Spark's advantage in Scala or SQL language. These conclusions are drawn mainly from the documentation of the two frameworks.

1. When to Spark

- When the use case is relatively simple, i.e. cleanly fits the Map-Shuffle-Reduce paradigm.
- When one can find a quick solution with the high-level tools provided by Spark.

2. When to use Dask

- When the case is complex and more customization is required.
- When one wants to perform parallelization to the existing legacy code

Existing documentations on comparison

- From Dask website: <https://docs.dask.org/en/latest/spark.html>
- Dask vs Spark comparison case on neuroimaging pipelines (concluding the performance is similar): <https://arxiv.org/pdf/1907.13030.pdf>